

Seed 2.0 Pro × Cowork-Bench

Past evaluation, present-day context, and meeting preparation

[READ THIS FIRST](#)

The report evaluated Seed 2.0 Pro, not Seedance

Keep the product names separate

Seed 2.0 Pro is ByteDance's general-purpose agent model. **Seedance** is its video-generation family. The report tested multi-step agent work inside Eigent; it did not evaluate video generation.

The one-sentence takeaway

Seed generally understood the tasks, but it was much less reliable at the final mile: exact schemas, efficient repeated operations, final artifact verification, and completing every requested deliverable.

0.22

Seed strict pass@3

1.00

Opus strict pass@3

0.77

Reported Seed lenient result

Why this work matters today

The benchmark surfaced reliability problems that matter in production agents, not merely benchmark intelligence. ByteDance's later Seed 2.1 release explicitly emphasizes reliable cross-tool execution and consistent delivery, making this report a useful historical baseline for today's discussion.

Meeting objective

Share Eigent's product direction, learn how their current agent and media-model products fit together, and explore a fair rerun of the same real-world benchmark on Seed 2.1.

PAST WORK

What we tested

The report was created on 5 May 2026. It tested whether Seed 2.0 Pro could complete realistic, multi-system knowledge-work tasks inside Eigent.

Scope	10 Cowork-Bench tasks × 3 attempts: 30 Seed trials and 30 Opus 4.7 trials
Work	GitHub operations, document creation, dashboards, data aggregation, Linear/Trello updates, deployments, and Slack publishing
Shared setup	The same 167 skills, 5,400-second timeout, isolated trial containers, and scoring process
Different harnesses	Seed ran through Eigent's CAMEL Workforce; Opus ran through Claude Code CLI

Headline results

Capability	Seed 2.0 Pro	Opus 4.7	Gap
Terminal use	0.59	0.98	−0.39
Search / data	0.56	1.00	−0.44

How to interpret the numbers

Strict scoring required the complete rubric. The lenient threshold accepted materially correct partial work. Pass@k means the chance of obtaining at least one acceptable result across k attempts; it is not single-run accuracy.

The fairest conclusion

This was not primarily a knowledge failure. Seed often found the right information and performed substantial work, but missed output details that production systems depend on.

Important evaluation caveat

Because the models used different agent harnesses, the benchmark measures the combined model-policy-runtime system. It should not be presented as a pure base-model comparison.

EVIDENCE

Where Seed lost reliability

1	Structured-output drift Invalid JSON, wrong field names, missing metadata, or only partial schema coverage.
2	Inefficient repeated operations Many individual CLI calls for homogeneous work instead of one script that loops over the items.
3	No final artifact verification The model wrote an output and moved on without rereading it against the user's exact request.
4	Workflow abandonment Some runs exhausted their budget before the final deliverable, or investigated the environment without starting the actual work.

Two memorable examples

Weather dashboard	Seed created valid JSON, then overwrote it with Python <i>repr()</i> output using single quotes. The grader could no longer parse the file.
Security review	Seed posted 15 PR replies in separate terminal calls, used incorrect JSON keys, and never produced the required Word document. Opus batched the replies in one script and completed both deliverables.

Report cautions to remember

- The benchmark is a small suite: ten tasks, three attempts each.
- One task was excluded from each model's strict pass@3 calculation because it had fewer than three valid attempts.
- The executive summary labels 0.77 as lenient pass@3, while the detailed table labels it pass@2 ≥ 0.5. In conversation, call it the **reported 77% lenient result** unless the source calculations are available.

Diplomatic framing

Present this as actionable external feedback and a historical baseline—not as a current verdict on Seed 2.1 and not as proof that every observed difference came from the base model.

TODAY'S MEETING

Current context and talking points

ByteDance released **Seed 2.1** on 23 June 2026, emphasizing reliable general-agent delivery, cross-tool workflows, and production coding. It released **Seedance 2.5** on 31 July 2026, focused on 30-second audiovisual storytelling, multimodal references, and precise editing.

Public Volcano Engine material identifies 吴迪 (Wu Di) as Head of Intelligent Algorithms and head of ModelArk/火山方舟. An authoritative “Head of Product” title was not found, but his public role is directly relevant to model productization and enterprise deployment.

A strong opening

Suggested wording

“We evaluated Seed 2.0 Pro on realistic multi-system workflows inside Eigent. The model generally understood the work, but the main gap was reliable completion: exact schemas, efficient repeated operations, and final artifact validation. We would love to understand how Seed 2.1 has addressed those areas and explore rerunning the same benchmark.”

Questions worth asking

1. Which Seed 2.1 improvements directly target long-horizon completion and cross-tool reliability?
2. What should be handled by the model versus the agent runtime—batching, retries, schema validation, and final verification?
3. What structured-output guarantees, observability, and tool-use capabilities are on their roadmap?
4. How do Seed, Seedance, Seedream, and ModelArk fit together commercially?
5. Would they support an apples-to-apples Seed 2.1 rerun inside Eigent, with shared failure traces and regression cases?

Your collaboration angle

Eigent can contribute realistic multi-service workflows, reproducible evaluation, and detailed execution traces. Their team can show how current models and platform features address the specific reliability failures. A controlled Seed 2.1 rerun gives both sides a concrete next step.

Useful sources

Original report: **Seed 2.0 Pro Capability Analysis on Cowork-Bench**, 5 May 2026

Seed 2.1: [official announcement](#)

Seedance 2.5: [official overview](#)

Wu Di: [Volcano Engine profile](#)